



Use Your Own CUDA ROS Node with NITROS

Swapnesh Wani | Dec 05, 2023





Agenda

- Overview
- NITROS
- Performance
- Managed NITROS APIs
- Demo examples
- Questions

Motivation

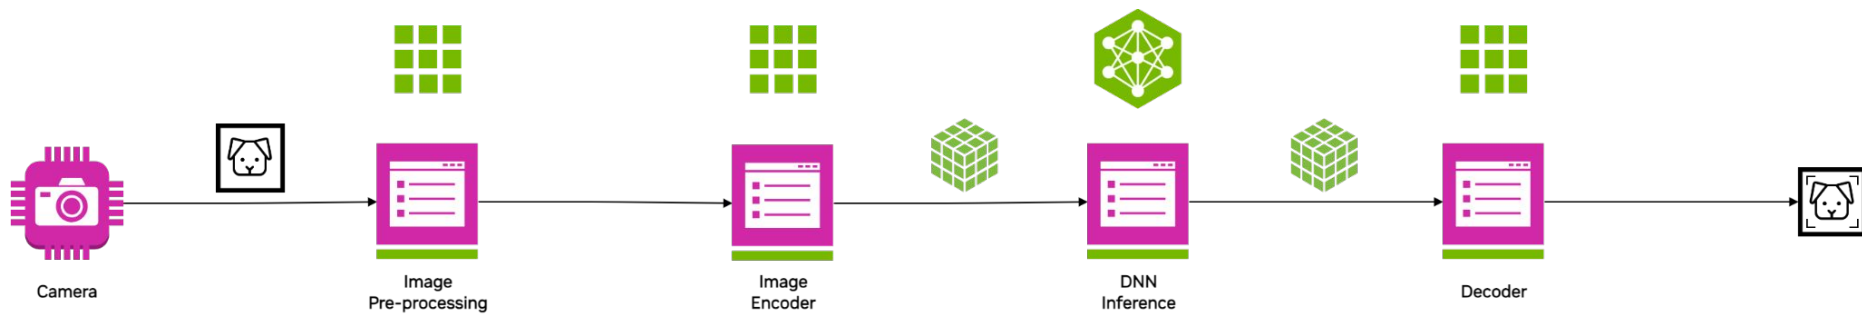
Make efficient use of CUDA in ROS2 nodes

- **Scenario** - You have an awesome CUDA code that is part of ROS 2 nodes and want to share the GPU buffer between them.
- **Problem** - Excessive memory copy between the two node.



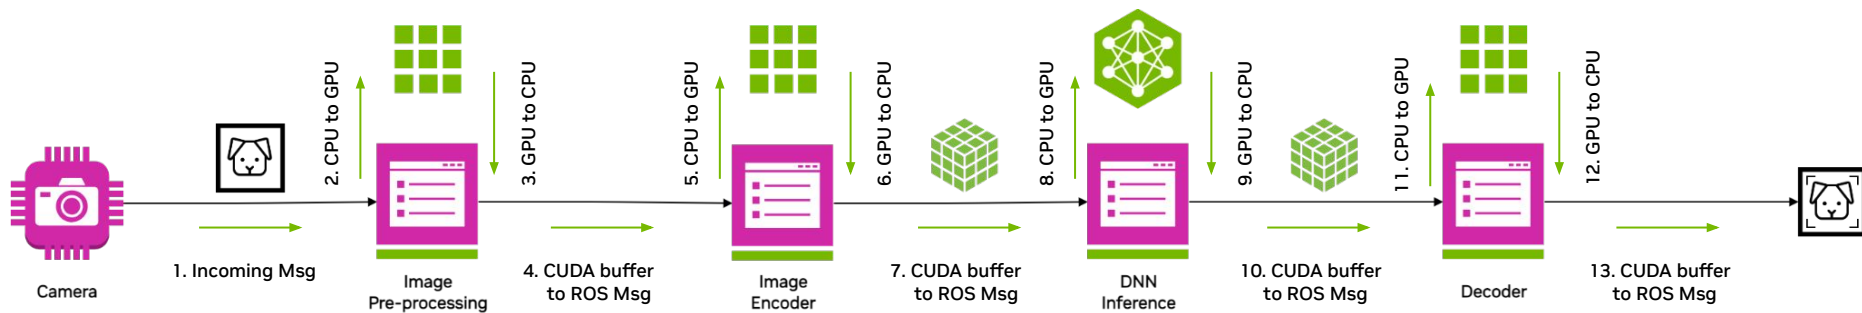
Real World Motivation

DNN inference pipeline



Real World Motivation

DNN inference details



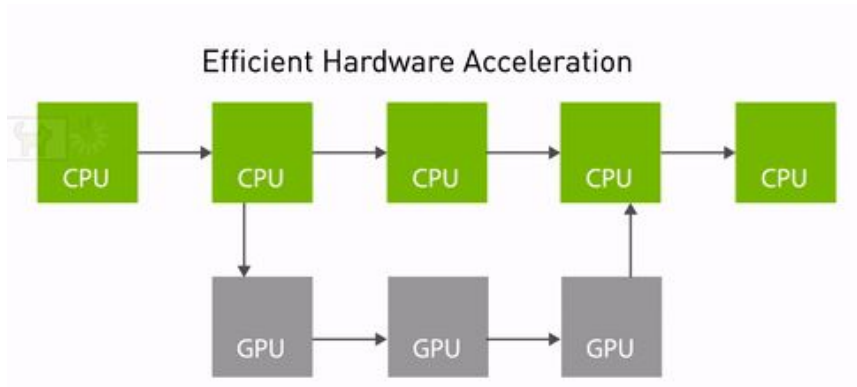
| Resolution | Image Size | Extra Mem Copy |
|------------|------------|----------------|
| 1280x720 | ~3 MB | ~18 MB |
| 1920x1080 | ~6.5 MB | ~39 MB |

6x penalty!

NITROS

Make efficient use of CUDA in ROS2 nodes

Solution



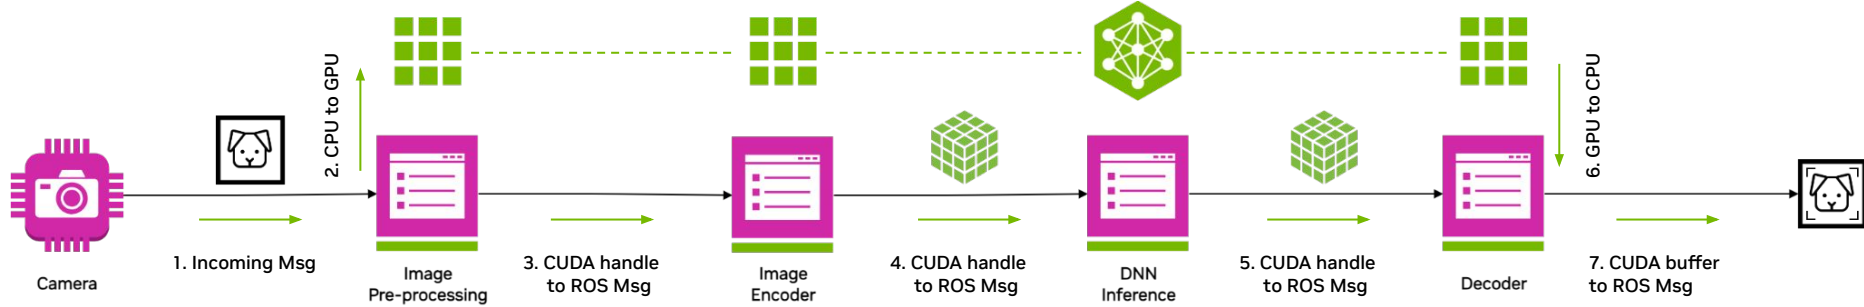
Isaac ROS NITROS

Benefits

- ✓ Efficient use of hardware
- ✓ Reduced memory copies
- ✓ Compatible with existing ROS 2 packages and tools
 - RViz / Foxglove
 - CLI tools like `ros2 topic echo <topic_name>`
 - Nav2, Moveit2, etc.

Real World Motivation

Now with NITROS



| Resolution | Image Size | Extra Mem Copy |
|------------|------------|----------------|
| 1280x720 | ~3 MB | 0 MB |
| 1920x1080 | ~6.5 MB | 0 MB |

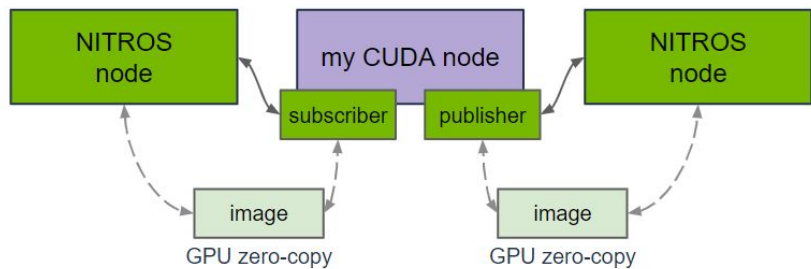
More **throughput** and less CPU load!

Performance

Of NITROS accelerated graphs

| Graph | Input Size | AGX Orin | AGX Xavier | Orin NX | Orin Nano 8GB | x86_64 w/ RTX 4060 Ti |
|-----------------------------------|------------|--------------------|--------------------|--------------------|--------------------|-----------------------|
| AprilTag Graph | 720p | 213 fps 14 ms | 122 fps 21 ms | 102 fps 20 ms | 71.0 fps 26 ms | 472 fps 9.5 ms |
| Freespace Segmentation Graph | 576p | 45.9 fps 41 ms | 19.2 fps 130 ms | 27.6 fps 95 ms | 21.3 fps 110 ms | 91.0 fps 30 ms |
| Centerpose Pose Estimation Graph | VGA | 36.1 fps 5.7 ms | 20.2 fps 16 ms | 19.4 fps 7.4 ms | 13.8 fps 12 ms | 50.2 fps 14 ms |
| DOPE Pose Estimation Graph | VGA | 39.8 fps 33 ms | 12.5 fps 160 ms | 17.3 fps 120 ms | - | 89.2 fps 15 ms |
| DNN Stereo Disparity Graph Full | 576p | 74.0 fps 20 ms | 21.0 fps 93 ms | 26.1 fps 42 ms | - | 191 fps 11 ms |
| DNN Stereo Disparity Graph Light | 288p | 260 fps 13 ms | 101 fps 20 ms | 116 fps 16 ms | - | 350 fps 12 ms |
| Stereo Disparity Graph | 1080p | 155 fps 15 ms | 84.7 fps 21 ms | 74.7 fps 21 ms | 49.0 fps 31 ms | 357 fps 3.0 ms |
| DetectNet Object Detection Graph | 544p | 232 fps 11 ms | 90.0 fps 19 ms | 105 fps 15 ms | 74.2 fps 22 ms | 644 fps 5.6 ms |
| TensorRT Graph PeopleSemSegNet | 544p | 421 fps 8.1 ms | 252 fps 12 ms | 238 fps 9.6 ms | 162 fps 13 ms | 704 fps 5.5 ms |

CUDA with NITROS



- ✓ Ability to share GPU buffer across NITROS-enabled Isaac ROS node without extra copy
- ✓ Enabled increased parallel compute between CPU and GPU
- ✓ Simple APIs to use the NITROS publishers and subscribers.

Managed NITROS Publisher

- Provides interface for publishing messages in NITROS-enabled ROS graph.
- API is comparable to standard `roscpp::Publisher`

```
// Declaring NITROS publisher

using namespace nvidia::isaac_ros::nitros;
std::shared_ptr<ManagedNitrosPublisher <
    NitrosTensorList>> nitros_pub_;

// Creating NITROS publisher

nitros_pub_ = std::make_shared<
    ManagedNitrosPublisher <
    NitrosTensorList>>
    (this, "my_topic_name",
    nitros_tensor_list_nchw_rgb_f32_t ::
    supported_type_name);
```

Managed NITROS Subscriber

- Provides interface for subscribing messages from NITROS enabled ROS graph.
- API is comparable to standard `roscpp::Subscriber`

```
// Declaring NITROS subscriber
```

```
using namespace nvidia::isaac_ros::nitros;  
  
std::shared_ptr<ManagedNitrosSubscriber <  
    NitrosTensorListView >> nitros_sub_;
```

```
// Creating NITROS subscriber
```

```
nitros_sub_ = std::make_shared<  
    ManagedNitrosSubscriber <  
        NitrosTensorListView >>(  
        this, "my_topic_name",  
        nitros_tensor_list_nchw_rgb_f32_t ::  
        supported_type_name,  
        std::bind(&MyDecoderNode::Callback,  
        This, std::placeholders::_1))
```

NITROS Builder

- Provides utility classes to create NITROS-typed messages.
- Classes offer builder-style interface which allows users to specify relevant fields during object creation one at a time.

```
// NITROS builder example
```

```
using namespace nvidia::isaac_ros::nitros;  
NitrosTensorList tensor_list =  
    NitrosTensorListBuilder()  
        .WithHeader(ros_header)  
        .AddTensor("tensor_1",  
                    foo_tensor)  
        .AddTensor("tensor_2",  
                    bar_tensor)  
        .Build();
```

NITROS Viewer

- Provides utility classes to access NITROS-typed messages.
- Classes offer methods to access relevant fields from a GPU backed buffer which can be copied back for further processing.

```
// NITROS viewer example
```

```
using namespace nvidia::isaac_ros::nitros;
```

```
void MyDecoderNode::Callback(const  
NitrosTensorListView & msg)
```

```
{
```

```
    auto tensor = msg.GetNamedTensor(tensor_name_);
```

```
    size_t buffer_size{tensor.GetTensorSize()};
```

```
    std::vector<float> results_vector{};
```

```
    results_vector.resize(buffer_size);
```

```
    cudaMemcpy(results_vector.data(),
```

```
                tensor.GetBuffer(), buffer_size,
```


```
                cudaMemcpyDefault);
```

```
}
```

DEMO - 1

DEMO - 1

DEMO - 2



Conclusion

What did we learn

- ✓ How to efficiently use CUDA and integrate it with existing NITROS-enabled ROS graph
- ✓ How excessive memory copy can introduce bubbles in a hardware pipeline and affects the throughput and how to avoid it
- ✓ NITROS APIs and how to use them



Questions?

Relevant links

[Nvidia Isaac ROS Documentation Home](#)

[Performance Summary Page](#)

[Isaac ROS NITROS](#)

[CUDA with NITROS](#)

[Isaac ROS Webinar Series](#)

[NVIDIA Developer Blogs](#)

[ROS Discourse](#)

